

Squawk Technology

Nik Shaylor
Doug Simon
Bill Bush



Sun Microsystems Laboratories

#1

Squawk Technology



Overview

- The Project
- Technical Issues



The Squawk Project



Sun Microsystems Laboratories

#3

Squawk Technology



The Squawk Project

- Technical Overview
- What we've done so far



The Squawk Project

- **Goal: CLDC/KVM compliant implementation for small devices**
- **Goal: system written in Java for portability and ease of development**



Small JVM Technology

- **Small Footprint**
 - The next generation of smart cards:
 - 8 K of RAM
 - 32 K of EEPROM
 - 160 K of ROM
 - 32 bit processor
 - Size issues include
 - Size of transmitted class files
 - RAM needed for class file loading
 - Size of loaded class files
 - RAM required during execution
 - Size of interpreter and memory system



Implemented in Java

- The whole VM is written in Java
 - Test: run under Hotspot
 - Deploy: translate to C and compiled using a C compiler
- The VM has two types of components
 - Core functionality (interpreter and GC)
 - System classes
 - Regular Java system classes
 - Loader
 - Verifier
 - Thread scheduler
 - Object synchronization



Split VM

- Off-Device Translator
 - Reads pre-verified class files
 - Writes space-optimized collections of classes called *Suites*
- Suites
 - Contain only inter-suite symbolic information
 - Internally linked
 - Compact class representation
 - Space-optimized bytecodes
 - Optimized to simplify verification and GC



Published Specification

- **Squawk Specification -- Draft 2.1**
 - Defines Architecture
 - Defines Suite file format
 - Defines Bytecode set
 - Defines Bytecode transformations for verification and GC
- **Specification has two forms**
 - A standard form in which all Java programs can be represented
 - A minimal form for small programs that limits quantities like the number of fields and methods in a class
- See <http://sunlabs.eng/projects/squawk>



Implementation Status

- A complete prototype implementation has been written
- Developed and tested in Java
- Translated to C and compiled on Solaris/x86/Mac



Conformance Tests

- Product-level testing is complete
 - All 4628 CLDC 1.0 TCK compatibility tests have been run
 - 4537 pass (98.05%)
 - 37 fail because VM constraints are exceeded (such as a class with more than 256 static fields)
 - 43 fail in the translator due to its current inability to handle esoteric or border case constructs
 - 11 fail due to the absence of complete runtime access control (the verifier currently ignores *private*, *protected*, and *public* access modifiers)



Results

- Results
 - Suites are between 35% and 45% size of J2ME class files
 - Less than 4K of RAM is needed to install small suites
 - The footprint of the core interpreter is about 20K (x86)
 - Execution speed is between 85% and 107% of KVM.

(These are all very early results)



Next Generation Java Card

- Collaboration with Java Card group
 - Artifact
 - Standards work
- Java Card Forum
 - Three contenders: JEFF, standard CLDC, Squawk
 - Choice will not be technical
- Ongoing work to adapt Squawk to the Java Card platform



Technical Issues



Sun Microsystems Laboratories

#14

Squawk Technology



Application architecture




Sun Microsystems Laboratories

#15

Squawk Technology



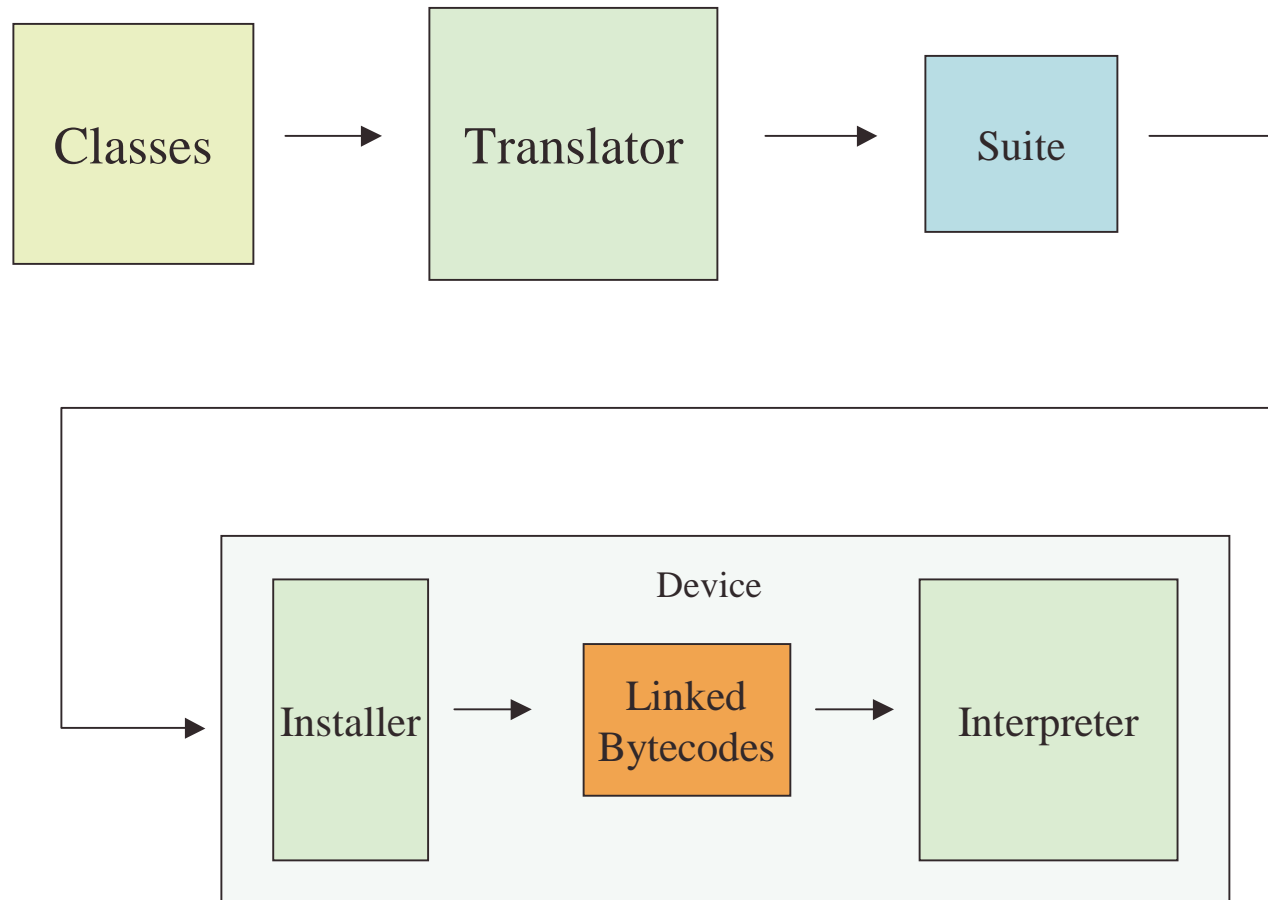
Suite construction



- Java source is converted into class files
- Class files are converted into a suite
- Suites are installed onto the Java card



Bytecode quickening



Translation



Sun Microsystems Laboratories

#18

Squawk Technology



Source example

A method from java.lang.Object

```
public boolean equals(Object obj) {  
    return (this == obj);  
}
```



javac output

```
Method boolean equals(java.lang.Object)
  0 aload_0
  1 aload_1
  2 if_acmpne 9
  5 iconst_1
  6 goto 10
  9 iconst_0
 10 ireturn
```



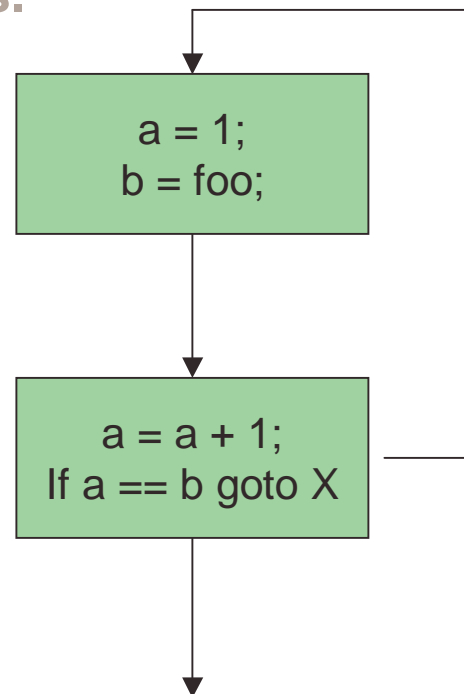
After translation (XML output)

```
<method_body>
  <type>1</type>
  <entry>7</entry>
  <locals>
    <type>1</type>
    <type>1</type>
  </locals>
  <stack>2</stack>
  <code>
    <load_0/>
    <load_1/>
    <if_icmpne/><byte>2</byte>
    <const_1/>
    <return/>
    <const_0/>
    <return/>
  </code>
</method_body>
```



Bytecode manipulation

- The Java stack must be empty at basic block boundaries.




This is an issue that simplifies verification.



Bytecode manipulation

- The stack must only contain the operands for certain operations such as invoke, getstatic etc.

Normal Java



Squawk



Evaluation Stack

This is an issue that greatly simplifies GC.



Bytecode manipulation

- Local variables can only be used for one data type.

Normal Java

Header
int/Object
float/byte[]

Squawk

Header
Object
byte[]
int
float

This is an issue that simplifies verification and GC.

Cost are minimal:

~5% increase in local variables.

~3% increase in code size.



Loading



Sun Microsystems Laboratories

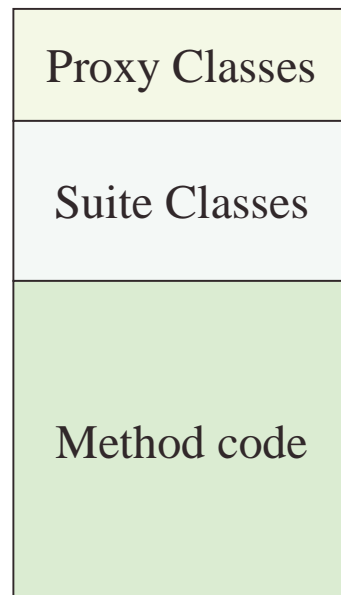
#25

Squawk Technology



Suites

- A suite is a complete collection of classes consisting of proxy classes, suite classes and method code.



Proxy Classes

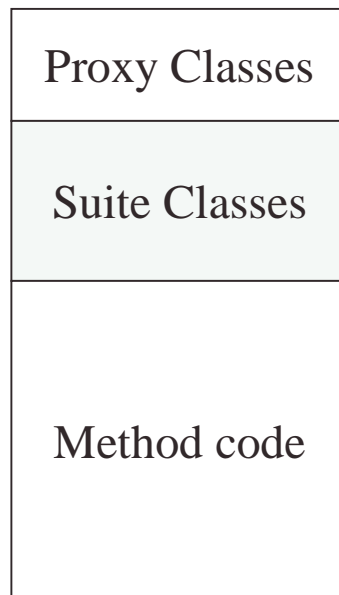
- Proxy classes represent classes that the suite will bind to at load time. They contain minimal symbolic information.

Proxy Classes
Suite Classes
Method code



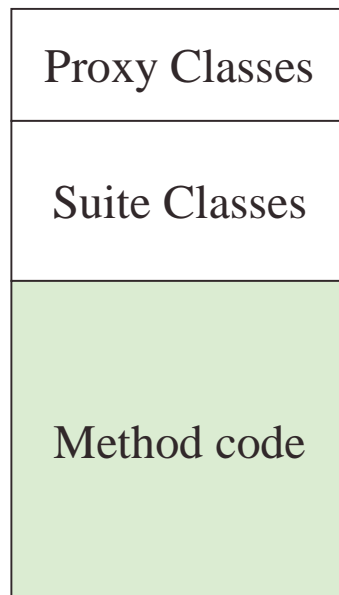
Suite Classes

- Suite classes contain the full information for a new class being loaded.



Method code

- The method code is the collection of bytecodes that define the methods of the suite classes.



Bytecode compression

- The Squawk bytecodes are smaller than regular Java bytecodes.

Opcode	Java	Squawk
goto	3 bytes	2 bytes
getfield	3 bytes	2 bytes
load_n	3 bytes	2 bytes
invokevirtual	3 bytes	2 bytes



Loading process

InputStream

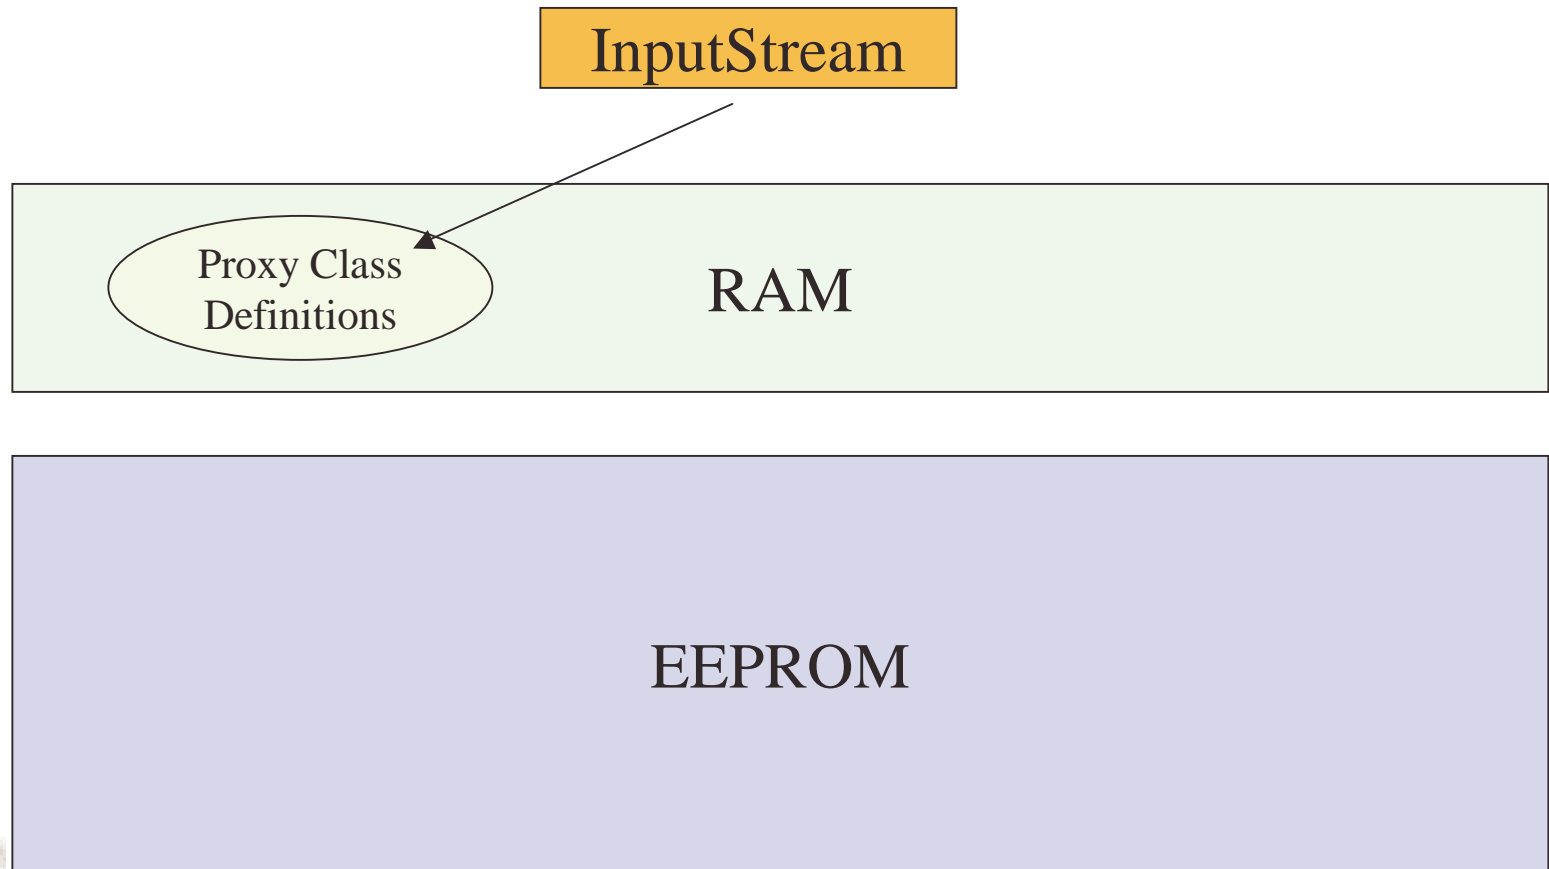
RAM

EEPROM

1 - An input stream is opened



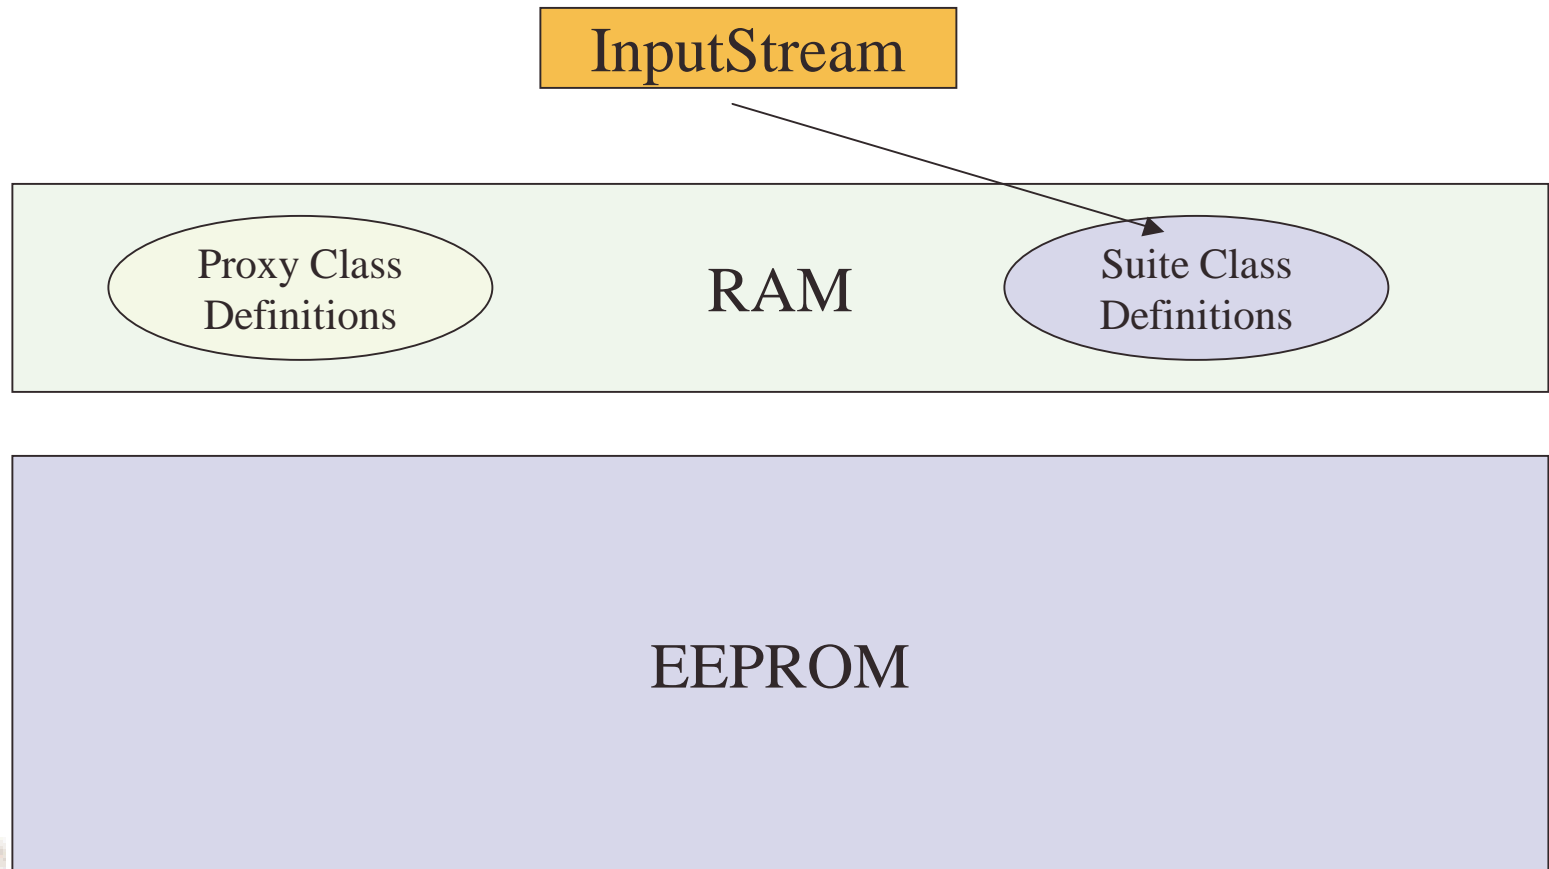
Loading process



2 – The Proxy class definitions are read into RAM



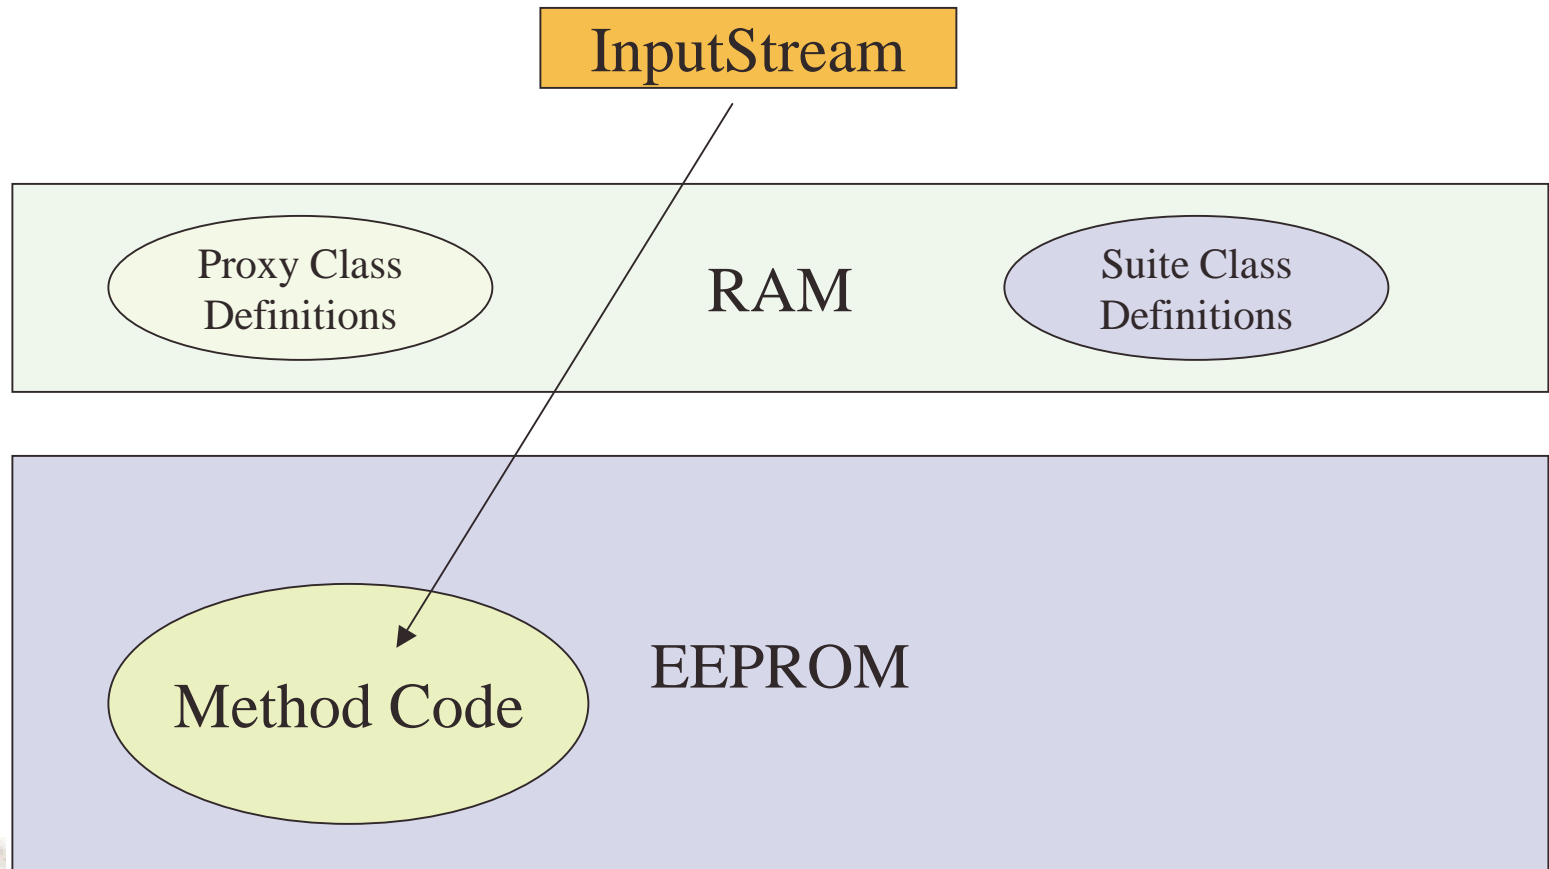
Loading process



3 – The Suite class definitions are read into RAM



Loading process



JAVA™

4 – The bytecodes are checked, converted, linked and written into EEPROM

Sun Microsystems Laboratories

#34

Squawk Technology



Loading process

InputStream

Proxy Class
Definitions

RAM

Suite Class
Definitions

Method Code

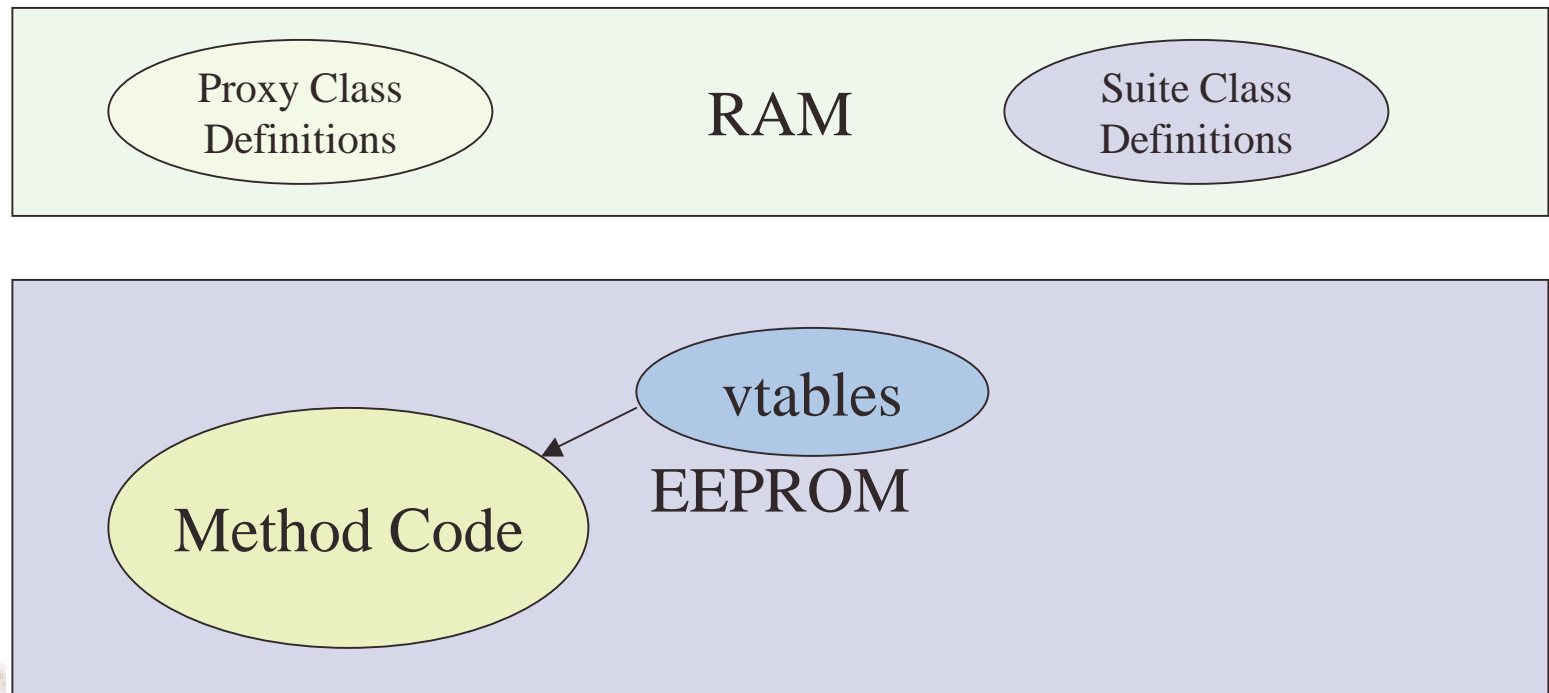
EEPROM

5 – A final check is made to see that the branch targets are all valid.



Loading process

InputStream

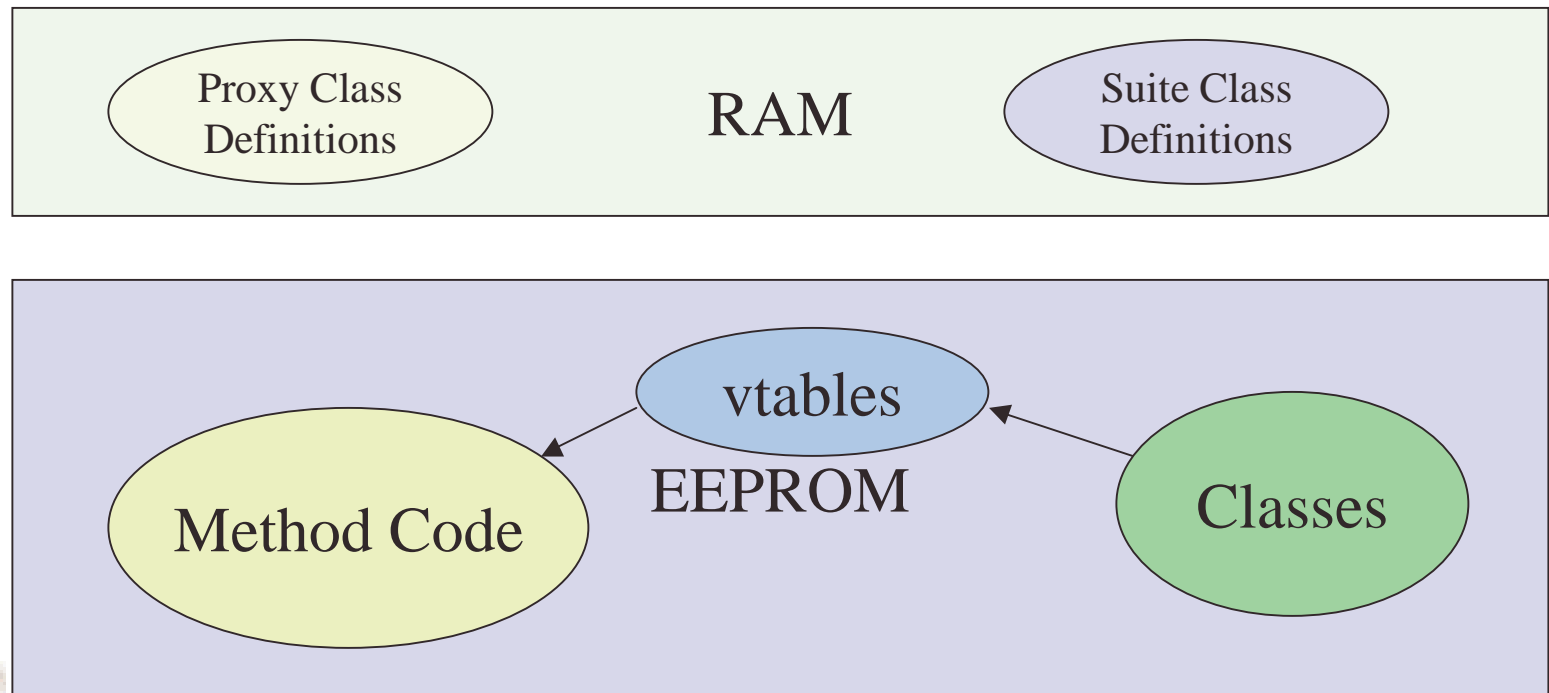


6 – The vtables for the methods are created



Loading process

InputStream

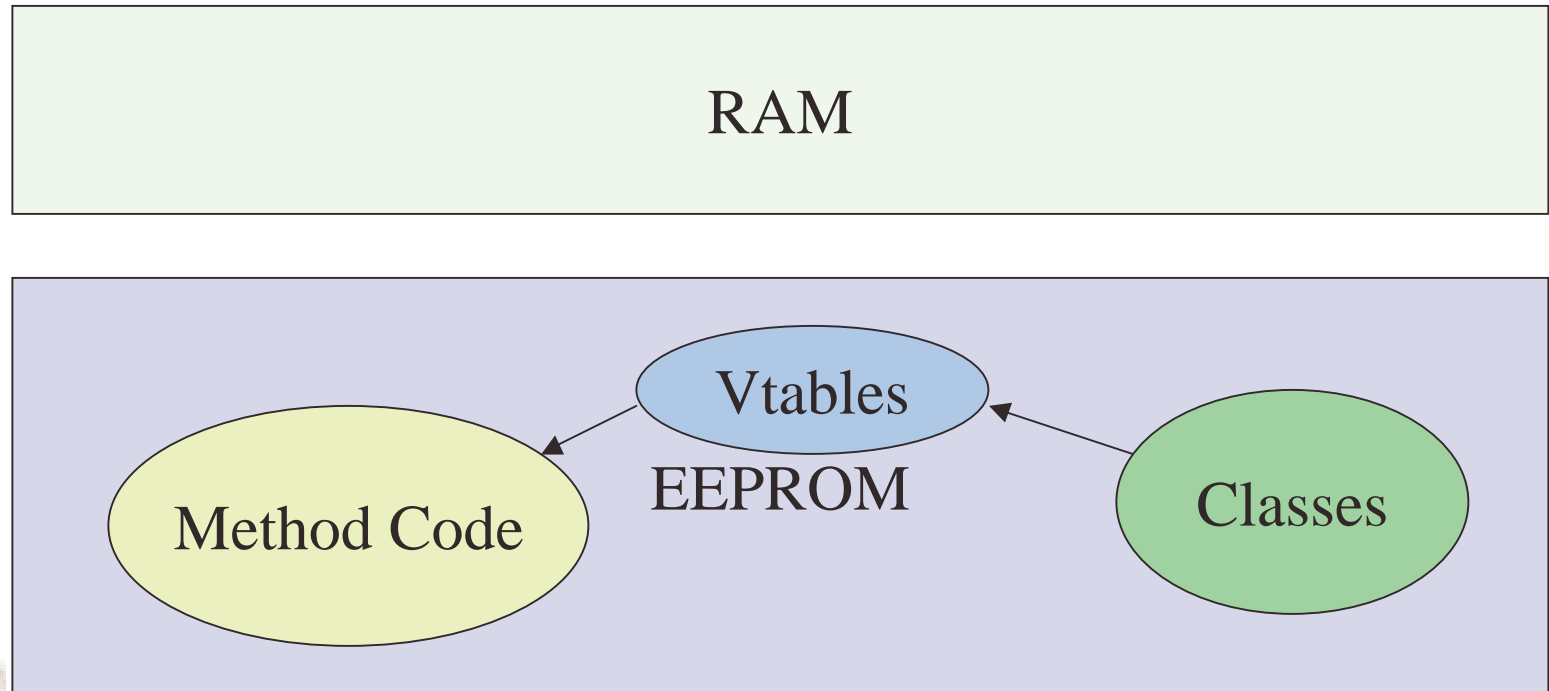


7 – The new class data structures are created



JAVA™

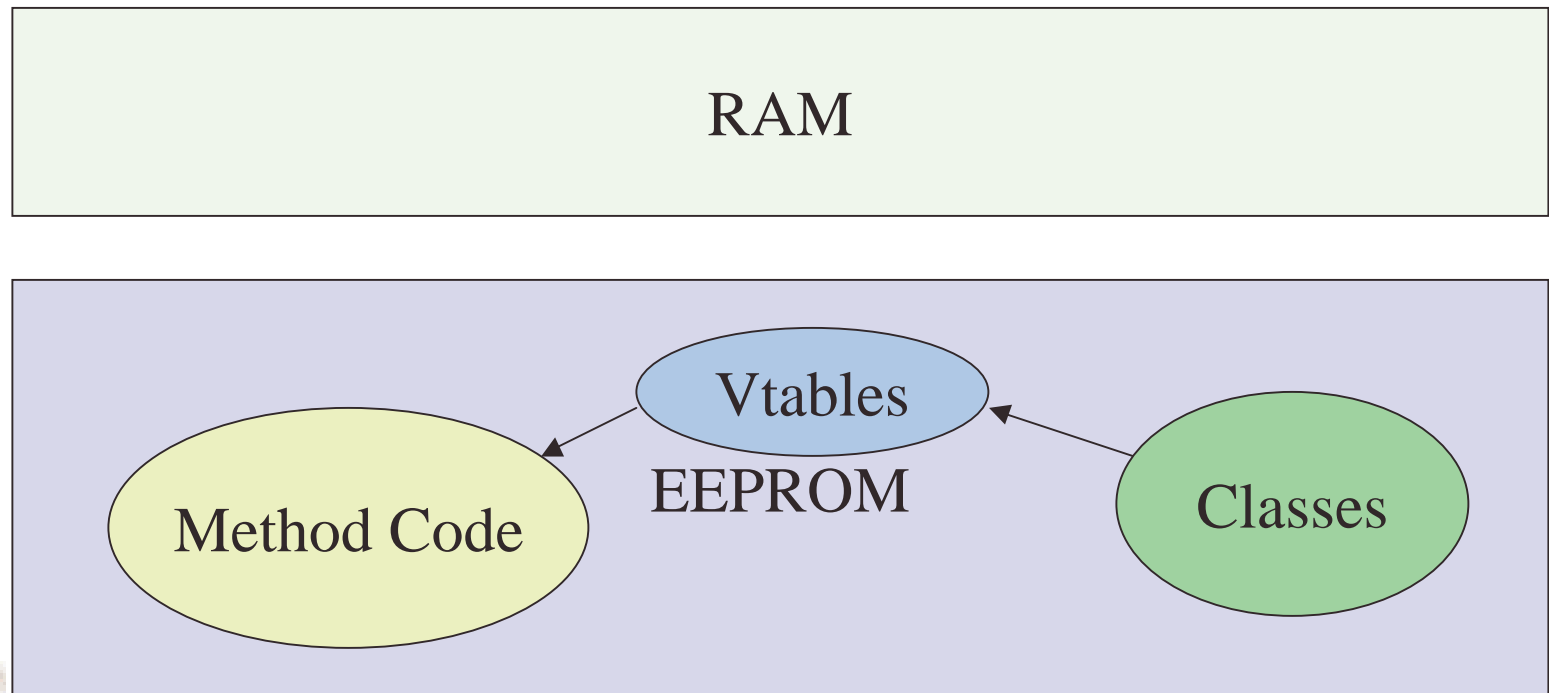
Loading process



8 – The temporary data structures are deleted



Loading process




Everything is now fully resolved, read-only, and ready for execution.



Memory model



Memory model




- Squawk uses three object memories.
- All Squawk data structures are Java objects.
- Different garbage collectors are used for the RAM and EEPROM.

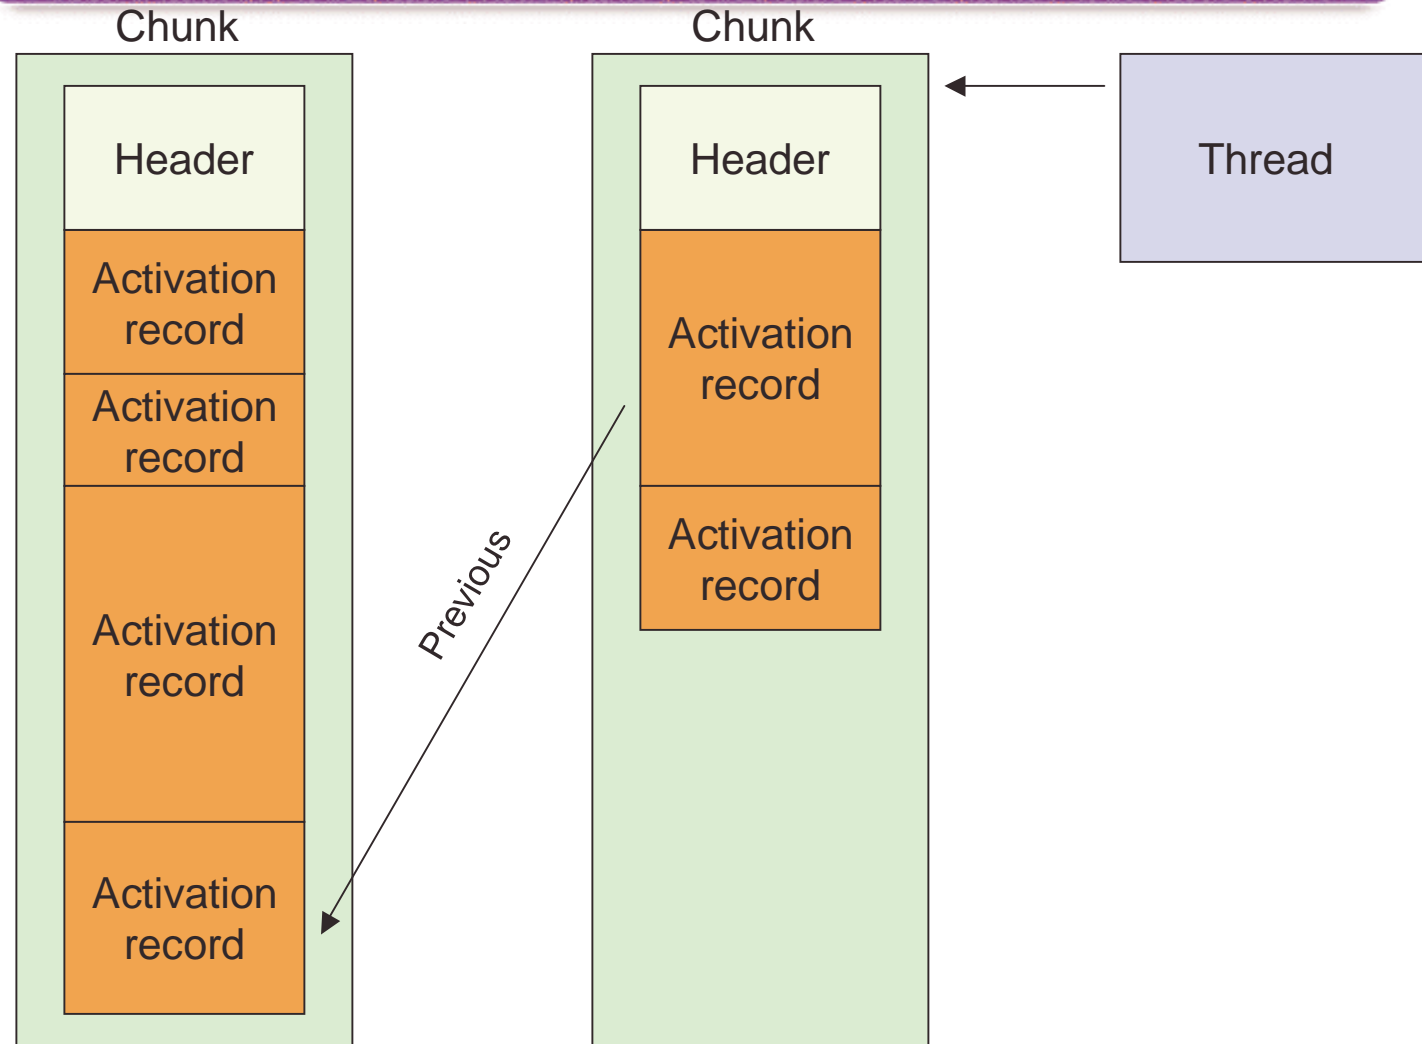


Class references

- Objects directly refer to the immutable information of their class. This can therefore be kept in ROM.
- The mutable class state is held separately in an associatively mapped location.



Activation records, Stack chunks, and Threads



JAVA™

Sun Microsystems Laboratories

#43

Squawk Technology



Garbage Collection



Sun Microsystems Laboratories

#44

Squawk Technology



GC Features


- All the collectors are exact.
- The translator makes exact garbage collection much easier
 - Local variables are either pointers or non-pointers
 - Sections of evaluation stack are not found between activation records.
- RAM Collector
 - Cheney collector
 - Two generational “Lisp 2” mark/sweep/compact collector
- EEPROM Collector
 - Non compacting mark/sweep.
 - Minimizes EEPROM writes (which are both slow & finite)
- Support for
 - Finalization
 - Object migration



Copying to EEPROM

```
Foo foo2 = (Foo)PersistentMemory.makePersistentCopy(foo1);
```

RAM




EEPROM




Copying to EEPROM

```
Foo foo2 = (Foo)PersistentMemory.makePersistentCopy(foo1);
```

RAM




EEPROM



Migrating to EEPROM

```
PersistentMemory.makePersistent(foo1);
```

RAM

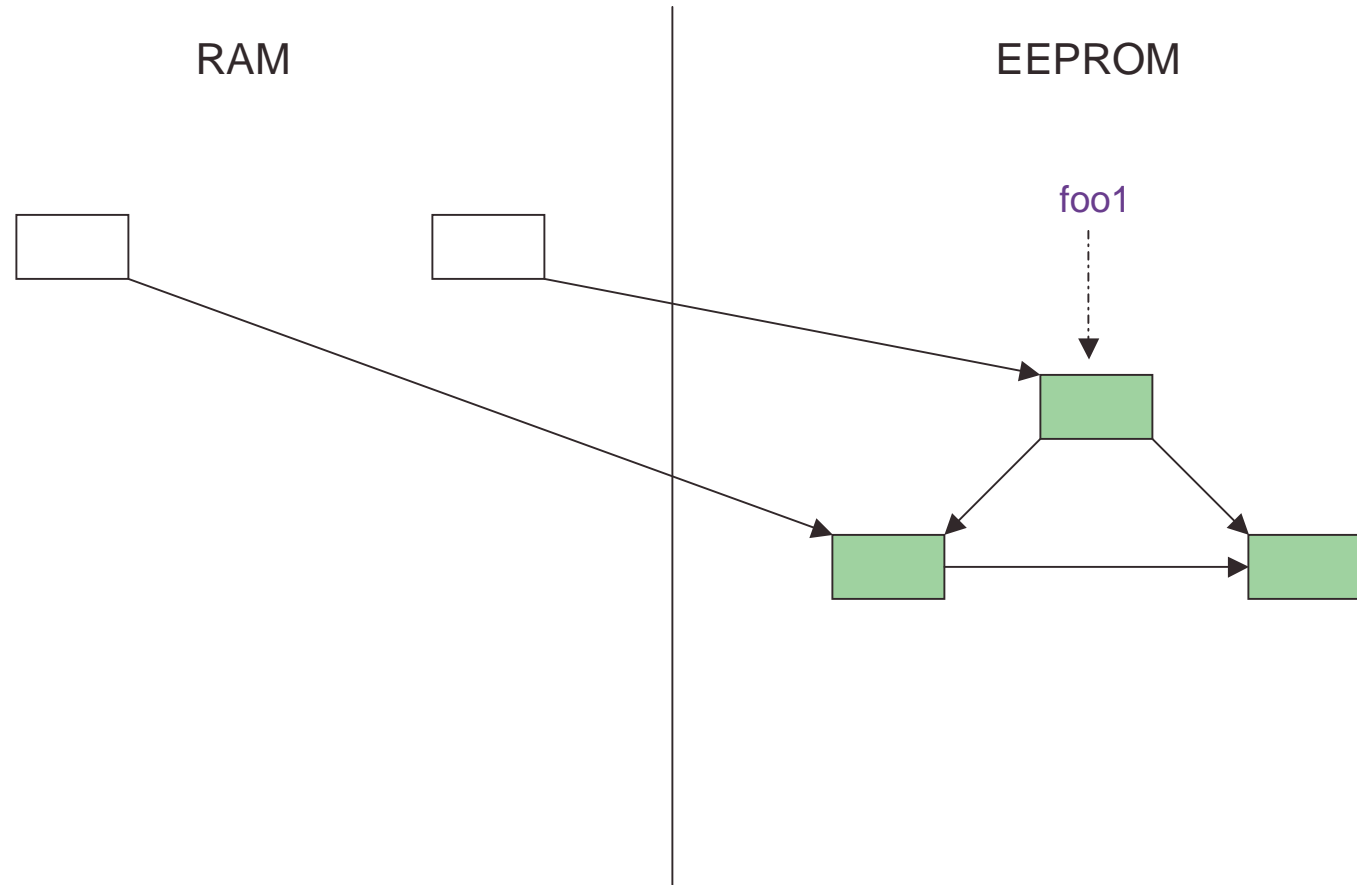


EEPROM



Migrating to EEPROM

```
PersistentMemory.makePersistent(foo1);
```



Interpreter core engineering



Sun Microsystems Laboratories

#50

Squawk Technology



Java to C conversion

All of the Squawk project is written in Java. The core interpreter is converted into C.

Interpreter.java



squawk.c

```
$ls -l *.java
```

```
17519 Dec 20 13:28 Interpreter.java
114021 Dec 18 17:30 Interpret.java
14871 Dec 18 17:30 Memory.java
99364 Dec 20 10:59 ObjectMemory.java
47358 Dec 23 14:12 PlatformAbstraction.java
```

```
$ls -l *.c
```

```
16972 Dec 23 14:39 squawk.c
113779 Dec 23 14:39 interp.c
14341 Dec 23 14:39 memory.c
93520 Dec 23 14:39 object.c
46381 Dec 23 14:39 platform.c
```



Example code

```
case OPC_IADD: { int r = pop() ; int l = pop() ; push(l + r); continue; }
case OPC_ISUB: { int r = pop() ; int l = pop() ; push(l - r); continue; }
case OPC_IAND: { int r = pop() ; int l = pop() ; push(l & r); continue; }
case OPC_IOR:  { int r = pop() ; int l = pop() ; push(l | r); continue; }
case OPC_IXOR: { int r = pop() ; int l = pop() ; push(l ^ r); continue; }
etc...
```

Most of the interpreter
is written in a subset of
Java and C.



Language differences

Original Java

```
void copyBytes(int src, int dst, int num) {
    if (num < 0) {
        fatalVMError("Negative range");
    }
    /*IFJ*/ System.arraycopy(memory, src, memory, dst, num);
    //IFC// memmove(memory+dst, memory+src, num);
}
```

Derived C

```
void copyBytes(int src, int dst, int num) {
    if (num < 0) {
        fatalVMError("Negative range");
    }
    /**** Line deleted by Squawk builder ****/
    memmove(memory+dst, memory+src, num);
}
```

Language differences are solved using a special form of conditional compilation



Feature elimination

```
/*if[FLOATS]*/  
  
    public final void writeFloat(float v) throws IOException {  
        writeInt(Float.floatToIntBits(v));  
    }  
  
/*end[FLOATS]*/
```

Features can be excluded from the interpreter core and the Java runtime libraries.



Native Methods



Sun Microsystems Laboratories


#55

Squawk Technology



Native methods

- Squawk supports a native method interface that can cause a Java thread to be suspended while it is waiting for an asynchronous event to take place.



Native methods

- However thread scheduling is never done in native code.
- Native methods are called using the following sequence:
 1. Execute native code
 2. Wait for event
 3. Get native result

```
public int read(byte b[], int off, int len) throws IOException {
    for (;;) {
        int result = execute(channel, b, off, len);
        if (result == 0) {
            return result(channel);
        }
        if (result < 0) {
            switch (result) {
                case EXNO_NoConnection: throw new ConnectionNotFoundException();
                case EXNO_IOException:  throw new IOException();
                default:                 fatalVMError();
            }
        } else {
            Thread.waitForEvent(result);
        }
    }
}
```



Native methods

- This interface has the following advantages:
 - The system is fully asynchronous allowing Java threads to execute while other threads are waiting for I/O.
 - Pointers into the object memory are never retained in native code because requests that cannot be satisfied immediately are always repeated.
 - This means that when the garbage collector runs there is never an object pointer in native code that needs updating.
 - Thread scheduling is all done in Java code.



The Demo



Sun Microsystems Laboratories

#59

Squawk Technology



Files


- `squawk.jar` contains the java version and the I/O library.
- `squawk.exe` is the C version.
- `squawk.image` is an image of the ROM, EEPROM, and RAM.
- `squawk.txt` is an description of how to use the demo.

```
36864 Jan  4 13:44 squawk.exe
4784160 Jan  4 13:37 squawk.image
204775 Jan  4 13:44 squawk.jar
3028 Dec 31 17:15 squawk.txt
```




I/O

- The version of the demo in C uses the same Java code for I/O.



Java version

Sun Microsystems Laboratories



C version

Squawk Technology

